


(a) Four architectures compared


Single network

369K · one perception, one decision




Separate experts, thin

363K · perception duplicated x5, each narrowed




Separate experts, full

1.83M · perception duplicated x5, full width




Shared perception (proposed)

512K · one perception, five specialised decisions




Radar perception is ~89% of single-network parameters.
The same expert decomposition is applied to MessageActor, ControlActor and Critic.


(b) Parameter count



(c) Proximity rate



(d) Arrival rate



v11 3 seeds per arm. Single-network (no MoE) is best on arrival and proximity; separate-full is worst. Shared hub sits in between.
This does not reproduce the YHSH story that shared MoE uniquely minimises crashes. Capacity vs specialisation still favours not duplicating the radar encoder (thin/full).